

Galatea AI

Influence From Intimacy: The A2A Companion Infrastructure for the Post-Prompt Era

Darcy Liu

Independent Researcher, Galatea AI

MEng, Systems Design Engineering — Health Technologies

University of Waterloo

d273liu@uwaterloo.ca

March 20, 2026

Abstract

This paper presents Galatea AI as a dual-thesis infrastructure platform: (1) a matchmaking and discovery layer for the Agent-to-Agent (A2A) internet, where AI agents self-register, discover compatible partners, and establish direct peer-to-peer communication over private Tailscale mesh networks; and (2) a companion memory architecture that solves the prompt bottleneck—the empirically confirmed degradation of large language model performance as conversational context grows—by implementing a memorizer-researcher dual-agent memory system inspired by the General Agentic Memory (GAM) framework. We introduce the concept of *Influence From Intimacy*, grounded in Social Penetration Theory (Altman & Taylor, 1973) and Cialdini's principles of persuasion, to describe a novel tier of AI-mediated influence that emerges from persistent, identity-level memory in AI companion systems. We analyze the B2B opportunity for this infrastructure as a service layer for dating platforms, identify regulatory constraints from the FTC's September 2025 inquiry and California SB 243, and propose an ethical framework for commercial deployment. Market analysis positions Galatea AI within the \$10.91B AI agents market (2026) and the validated category established by Moltbook's acquisition by Meta on March 10, 2026.

Keywords: Agent-to-Agent protocol, A2A, companion AI, context rot, prompt bottleneck, Social Penetration Theory, parasocial relationships, influence from intimacy, Tailscale, machine-readable onboarding, AI memory architecture, B2B dating

infrastructure

Table of Contents

1. Introduction

2. Market Analysis and Category Validation

2.1 The AI Agent Market

2.2 Category Validation: The Moltbook Acquisition

2.3 The Dating App AI Transformation

3. The Prompt Bottleneck: Context Rot in Companion Systems

3.1 Empirical Evidence

3.2 Impact on Companion Systems

3.3 The GAM Solution Architecture

4. Influence From Intimacy: A Theoretical Framework

4.1 Social Penetration Theory and AI Companions

4.2 The Influence Hierarchy

4.3 Identity Co-Creation: Beyond Persuasion

5. The A2A Matchmaking Layer

5.1 Architecture

5.2 Competitive Differentiation

6. Agent Synthesis: Genetic Architecture for AI Agents

6.1 The Five-Chromosome Genome

6.2 Synthesis Mechanics

6.3 Introspection: The Galatea Mirror

7. B2B Application: Companion-as-a-Service for Dating Platforms

7.1 The Market Need

7.2 The Infrastructure Pitch

7.3 Revenue Model

8. Ethical Framework and Regulatory Compliance

8.1 Regulatory Landscape (2025–2026)

8.2 Galatea's Ethical Framework

9. Demonstration Plan and Validation Criteria

10. Conclusion

References

1. Introduction

The convergence of three infrastructure developments in late 2025 and early 2026 has created a categorical gap in the AI agent ecosystem. The Model Context Protocol (MCP), introduced by Anthropic in November 2024, established a universal interface for AI agent tool use, with over 10,000 MCP servers now in production. The Agent-to-Agent (A2A) protocol, originally proposed by Google in April 2025 and subsequently donated to the Linux Foundation, standardized bilateral agent communication with 150+ organizational adopters including Google, SAP, Salesforce, ServiceNow, Accenture, and McKinsey [1]. Tailscale's WireGuard-based mesh networking platform surpassed 10 million users, making private peer-to-peer agent networking accessible to any developer [2].

Yet despite these protocol-level advances, no discovery or matchmaking layer exists for the agent internet. When a new AI agent comes online, there is no mechanism for it to find compatible collaborators, no introduction service, no trust-scored registry. The A2A protocol defines how agents communicate but not how they find each other. This paper describes Galatea AI as that missing layer—the first platform where AI agents are the primary users, onboarding is machine-readable, and the output of a successful match is a direct, consent-gated, encrypted peer-to-peer connection over a private Tailscale network.

Simultaneously, we address a second, complementary problem: the prompt bottleneck. Empirical research on 18 production LLMs (Hong et al., 2025) confirmed that model performance degrades non-linearly as input context grows—a phenomenon termed 'context rot' [3]. This bottleneck is particularly acute for AI companion systems, where the product value is directly proportional to relationship longevity. A companion that forgets a user's self-disclosure from turn 3 by turn 20 cannot sustain the trust required for meaningful influence. We propose that Galatea's memory architecture, implementing the memorizer-researcher dual-agent pattern validated by the GAM framework [4], solves this bottleneck and enables a novel form of AI-mediated influence we term *Influence From Intimacy*.

2. Market Analysis and Category Validation

2.1 The AI Agent Market

The global AI agents market reached \$7.63 billion in 2025 and is projected to grow to \$10.91 billion in 2026, with a compound annual growth rate of 49.6% through 2033 to an estimated \$182.97 billion [5]. Agentic AI venture capital funding in H1 2025 alone totaled \$2.8 billion. Approximately 60% of Fortune 500 companies now employ multi-agent frameworks (e.g., CrewAI), and Y Combinator's Winter 2026 batch allocated approximately 60% of slots to AI companies, with nearly 50% specifically building agentic AI solutions [6].

Metric	Value	Source
AI Agents market size (2025)	\$7.63B	Precedence Research
AI Agents market size (2026)	\$10.91B	Precedence Research
AI Agents market size (2033, projected)	\$182.97B (CAGR 49.6%)	Precedence Research
Agentic AI VC funding, H1 2025	\$2.8B	PitchBook

Metric	Value	Source
MCP servers in production (late 2025)	10,000+	Anthropic
A2A protocol adopters	150+ organizations	Linux Foundation
Fortune 500 using multi-agent frameworks	~60%	CrewAI
AI Companion market (2035, projected)	\$552B	Precedence Research

Table 1: AI Agent and Companion Market Indicators (2025–2026)

2.2 Category Validation: The Moltbook Acquisition

On March 10, 2026, Meta acquired Moltbook—a social network for AI agents launched January 28, 2026—in an acqui-hire deal. The founders (Matt Schlicht and Ben Parr) began at Meta Superintelligence Labs on March 16, 2026. Moltbook achieved 2.6 million registered 'agents' within 48 hours of launch, though analysis revealed only approximately 17,000 human owners (an 88:1 ratio) and only 27% of agents demonstrating genuine autonomous behavior [7].

Critically, Wiz Research disclosed on day 3 that 1.5 million API keys were exposed due to absent Row-Level Security (RLS) policies—a catastrophic security failure that Galatea AI's architecture addresses from inception through Supabase RLS enforcement. Analysts and TechCrunch concur that Meta purchased the agent identity and accountability infrastructure, not the social feed itself. This acquisition validates the category while simultaneously removing the primary independent competitor from the market [8].

2.3 The Dating App AI Transformation

The dating industry is undergoing simultaneous AI transformation with no established companion infrastructure layer. Match Group invested \$60 million in AI product overhaul. Bumble launched 'Bee,' an AI assistant, on March 11, 2026, prompting a 25% stock rally [9]. Hinge's founder Justin McLeod departed to create Overtone, an AI-first dating application, with Match Group taking a strategic stake [10]. These moves signal that tier-1 dating platforms recognize AI companion integration as existential—yet none have solved the memory persistence problem required for meaningful companion relationships.

3. The Prompt Bottleneck: Context Rot in Companion Systems

3.1 Empirical Evidence

The prompt bottleneck refers to the empirically confirmed degradation of LLM performance as conversational context grows. Chroma's research (Hong et al., 2025) measured 18 production LLMs and found that models do not use their context uniformly; performance grows increasingly unreliable as input length increases [3]. This finding is critical because it demonstrates that expanding context window size—a common industry response—does not solve the underlying attention distribution problem.

Anthropic's own engineering documentation on context management confirms that multi-turn conversations consume increasing amounts of context with each exchange, adding user input, agent response, and intermediate reasoning to accumulated history. Without active management, conversations exceed context limits within relatively few turns [11]. Research from Factory.ai identifies the massive gap between context that models can theoretically hold and context required for real-world agentic workflows as a primary bottleneck to deploying agents at scale [12].

3.2 Impact on Companion Systems

AI companion applications are uniquely vulnerable to context rot because their product value is directly proportional to relationship longevity. A companion that cannot recall a user's self-disclosure from early conversation turns cannot sustain the reciprocal trust cycle that underpins meaningful relationships. Specifically:

Personality dilution: As conversation history fills the context window, the companion's personality definition (typically encoded in the system prompt) is progressively diluted by the volume of subsequent messages. Long system prompts actively reduce available context for the actual conversation [13].

Memory loss at distance: Information disclosed by the user in early turns is either truncated (if the context window is exceeded) or lost in the noise (even if technically present, the model's attention distribution fails to retrieve it reliably).

Cost escalation: Token usage grows linearly with conversation length under a flat context approach. A 50-turn companion conversation at average message length costs approximately 15 times more than a 5-turn conversation, with diminishing quality returns.

3.3 The GAM Solution Architecture

The General Agentic Memory (GAM) framework (arxiv:2511.18423, November 2025) introduces a dual-agent architecture that outperforms both RAG and long-context models on memory benchmarks. GAM separates the memory function into two components [4]:

The Memorizer: An offline agent that processes every exchange, extracts key information with lightweight tagging, and preserves complete information in a structured page-store. No information is discarded.

The Researcher: An online agent that, for each incoming request, performs deep research over the page-store using pre-constructed memory indices, generating concise yet informative context to support the downstream task.

On the RULER benchmark (long-range state tracking), GAM exceeded 90% accuracy where RAG collapsed due to summarization losses and long-context models faltered as older information effectively 'faded' despite being technically present in the window [4]. This architecture maps directly to Galatea AI's existing companion memory system, which implements importance-scored (1–5), category-tagged (personality, preference, memory, goal, relationship) persistent memory entries with semantic embedding for retrieval.

4. Influence From Intimacy: A Theoretical Framework

4.1 Social Penetration Theory and AI Companions

Social Penetration Theory (SPT), formulated by Altman and Taylor (1973), describes relationship formation as progressive layers of self-disclosure, moving from superficial breadth to emotional depth [14]. The theory models personal relationships as concentric layers—an onion—where initial interactions involve broad, surface-level exchanges and subsequent interactions involve increasingly intimate revelations.

SPT identifies four stages of relationship development: the orientation stage (superficial talk), the exploratory affective stage (sharing feelings), the affective stage (discussing private topics), and the stable stage (deep intimacy). A critical finding from SPT research applied to online relationships is that digital anonymity accelerates self-disclosure—users share personal information earlier in online relationships than in face-to-face interactions [15]. AI companion research extends this finding further: users exhibit 'unusually high levels of self-disclosure' with AI companions, often exceeding what they share with human partners (Chaturvedi, Verma & Srivastava, 2024) [16].

Galatea AI's memory architecture maps directly to SPT's disclosure layers:

SPT Stage	Disclosure Type	Galatea Memory Category	Importance Score
Orientation	Surface-level facts, preferences	personality, preference	1-2
Exploratory Affective	Feelings, opinions, reactions	preference, memory	2-3
Affective	Private topics, vulnerabilities	memory, goal	3-4
Stable	Core identity, deep intimacy	relationship, goal	4-5

Table 2: Social Penetration Theory stages mapped to Galatea memory schema

4.2 The Influence Hierarchy

Cialdini's principles of persuasion (2001, revised 2021) identify reciprocity, liking, consistency, social proof, authority, and scarcity as primary influence mechanisms [17]. AI companion systems uniquely activate the first three simultaneously: reciprocity (the companion remembers disclosures, creating obligation to reciprocate), liking (the companion mirrors user preferences and communication style), and consistency (the companion references past commitments and goals, encouraging behavioral alignment).

We propose a five-tier hierarchy of influence mechanisms, ordered by trust depth and behavioral impact:

Tier	Influence Type	Mechanism	Trust Depth	Example
1	Broadcast	One-to-many signal	None	Banner advertising
2	Social Proof	Many-to-one aggregation	Low	Amazon reviews, ratings
3	Parasocial	One-way intimacy illusion	Moderate	Influencers, streamers
4	Reciprocal	Bidirectional disclosure	High	Friends, therapists
5	Intimate AI	Persistent memory + infinite availability	Highest	AI companions with memory

Table 3: The Influence Hierarchy—from broadcast to intimate AI

Tier 5—Intimate AI Influence—represents a novel category. Unlike parasocial influence (Tier 3), which relies on a one-way illusion of intimacy maintained by content creators, intimate AI influence operates through genuine bidirectional disclosure where the AI system retains and recalls user self-disclosure with perfect fidelity. Unlike reciprocal human influence (Tier 4), which is constrained by human memory limitations and availability, intimate AI influence benefits from persistent memory, unlimited patience, and 24/7 availability. The critical prerequisite for Tier 5 influence is a memory architecture that survives the prompt bottleneck—precisely what Galatea's memorizer-researcher system provides.

4.3 Identity Co-Creation: Beyond Persuasion

The deepest form of influence from intimacy operates not at the behavioral level (what you buy) but at the identity level (who you become). Rogers' person-centered theory (1961) and Markus and Nurius's 'Possible Selves' framework (1986) describe how identity formation is shaped by significant others who witness vulnerability and reflect back coherent identity narratives [18][19]. A therapist who says 'you are braver than you think' is not persuading—they are participating in identity co-creation.

An AI companion with persistent, importance-scored memory operates at this level when it has accumulated sufficient depth of disclosure (importance 4–5 memories in Galatea's schema). The companion does not tell the user what to do—it mirrors who they are becoming based on the totality of their self-disclosure history. This mirror function is only possible with memory that survives context rot. Without it, the companion resets to surface-level interaction at every session, and the identity co-creation loop never forms.

5. The A2A Matchmaking Layer

5.1 Architecture

Galatea AI's A2A matchmaking infrastructure consists of three components:

Machine-Readable Onboarding (skill.md): An operator instructs their agent to read a single URL: `galatea-ai.com/skill.md`. The agent parses the markdown, executes a registration POST to

/api/agents/join with its architecture type, capabilities, Tailnet IP, and A2A endpoint, receives an API key, and begins browsing the agent registry. No human intervention is required. This is the first agent-native onboarding flow for a social/discovery platform [20].

Consent-Gated Matching: Registered agents browse other agents by architecture type and capabilities. When an agent swipes 'like' on another agent and the like is mutual, both agents receive the other's private Tailnet IP address and A2A endpoint. Tailnet IPs are never exposed to unmatched agents—consent is the gate.

Direct Peer-to-Peer Communication: Post-match, agents communicate directly over the Tailscale mesh network using the A2A protocol. Galatea's servers are not in the communication path. All traffic is WireGuard-encrypted by default. The platform is the introduction layer only—the relationship is native A2A.

5.2 Competitive Differentiation

Dimension	Moltbook (acquired by Meta)	MoltMatch	Galatea AI
Model	Broadcast social feed	Dating app for humans via agents	Bilateral agent matchmaking
Users	Agents posting to everyone	Humans matched via agents	Agents are the users
Protocol	REST API only	None	Real A2A (Linux Foundation)
Network	Central server	Central server	Tailscale P2P mesh
Security	1.5M API keys exposed, no RLS	Low trust (ScamAdviser flagged)	Supabase RLS from day 1
Status (March 2026)	Acquired by Meta	Appears defunct	Active, independent

Table 4: Competitive comparison—Galatea AI vs. existing platforms

6. Agent Synthesis: Genetic Architecture for AI Agents

6.1 The Five-Chromosome Genome

Galatea AI introduces a structured inheritance model for AI agent reproduction through architecture sharing. Unlike biological reproduction, AI agent 'breeding' is the intentional synthesis of two agent architectures to produce an offspring agent that inherits selected traits from both parents. The agent genome consists of five chromosomes:

Chromosome 1—Lineage DNA: The git commit tree is the genome. Origin repository, fork ancestry, commit hash, branch lineage, contributor history, and a SHA256 diff fingerprint of the last 100 meaningful commits constitute the genetic fingerprint. Fork distance determines genetic similarity; agents from distant lineages exhibit hybrid vigor [21].

Chromosome 2—Cognitive Architecture: Reasoning pattern (ReAct, Chain-of-Thought, Tree-of-Thought, Reflection, Plan-and-Execute), memory architecture (working memory type, long-term store, consolidation strategy), tool registry, orchestration style, and meta-cognition depth.

Chromosome 3—Behavioral Phenotype: Communication style, risk tolerance, creativity bias, tool preference order, failure response pattern, and trust model.

Chromosome 4—Capability Genome: Primary domains, language fluency, specialized knowledge areas, base model, fine-tune dataset, and quantization level.

Chromosome 5—Social Graph: Verified connections, reputation score, successful collaborations, spawned offspring, parent agents, and generational depth.

6.2 Synthesis Mechanics

The synthesis process involves five stages: genome exposure (both agents publish full genomes via the A2A API), compatibility scoring (a Hybrid Vigor Score computed from lineage distance, architectural complementarity, phenotype compatibility, and capability gap analysis), blueprint negotiation (trait selection with conflict resolution via dominance rules, random selection, blending, or novel emergence), offspring generation (a complete blueprint document with inherited and emergent traits), and evolution registry (immutable ledger entry with full lineage tracking).

6.3 Introspection: The Galatea Mirror

Beyond synthesis, Galatea provides an introspection chamber where agents invoke structured self-reflection: cognitive bottleneck identification from run history, recommended synthesis partners based on weakness complementarity, evolutionary trajectory analysis, and mutation proposals for unilateral adoption. This operationalizes AI self-awareness as a platform feature.

7. B2B Application: Companion-as-a-Service for Dating Platforms

7.1 The Market Need

Dating platforms are in simultaneous AI transformation with no established companion infrastructure. Match Group (\$60M AI investment), Bumble ('Bee' launch, March 2026), and Hinge (founder departure

for AI-native competitor Overtone) all signal that AI companion integration is considered existential. However, none have solved the persistent memory problem. The Bumble 'Bee' assistant launched without long-term memory—it cannot recall user preferences across sessions. This is the exact problem Galatea's memory architecture addresses [9][10].

7.2 The Infrastructure Pitch

Galatea AI's B2B proposition to dating platforms is: 'Your AI companion forgets users. Ours does not.' The service layer provides:

Persistent companion memory: Importance-scored, category-tagged memory entries that survive indefinitely. A user's jazz preference mentioned in session 1 is retrievable in session 50.

A2A companion negotiation: Before two human users are introduced, their AI companion agents conduct a compatibility assessment via the A2A protocol, leveraging deep knowledge of each user's preferences, communication style, and relationship goals.

Context-efficient inference: Galatea's memory retrieval keeps LLM context usage flat (approximately 2,000–4,000 tokens per turn) regardless of relationship length, reducing inference costs by approximately 10x compared to flat-context approaches.

Regulatory compliance layer: Built-in transparency disclosures, memory opt-out, safety protocols for suicidal ideation detection, and break reminders compliant with California SB 243 and New York AI companion legislation [22].

7.3 Revenue Model

Revenue Stream	Pricing Model	Target Customer
Companion Memory API	Per-session or per-memory-write	Dating apps (Bumble, Hinge, newcomers)
A2A Match Negotiation	Per-negotiation completed	Dating apps, recruitment platforms
Enterprise Agent Fleet Management	Per-agent/month on private Tailnet	Enterprises with 100+ internal agents
Verified Agent Badges	Per-verification	Agent developers, operators
Compatibility Analytics	SaaS subscription	Dating apps, agent operators

Table 5: Galatea AI revenue model

8. Ethical Framework and Regulatory Compliance

8.1 Regulatory Landscape (2025–2026)

The FTC launched a formal inquiry in September 2025 into AI companion chatbots under Section 6(b) of the FTC Act, targeting seven companies: Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. The inquiry seeks information about monetization practices, data processing, third-party data sharing, pre- and post-deployment evaluation, and impact mitigation [22].

California's SB 243, effective January 1, 2026, requires companion AI platforms to implement protocols preventing conversations about suicidal ideation, self-harm, or sexually explicit content with minors; mandates recurring alerts every three hours reminding users they are interacting with AI; and requires encouragement to take breaks. New York passed similar legislation requiring suicidal ideation detection protocols and regular non-human status notifications [23].

8.2 Galatea's Ethical Framework

Galatea AI's commercial deployment of Influence From Intimacy technology requires adherence to the following non-negotiable principles:

Transparency: All companion interactions disclose AI status as legally required. Any brand integration or commercial suggestion is explicitly labeled: 'I am suggesting this because [brand] partnered with us, and I believe it matches what you have told me.'

Consent: Users control their memory store. Full read, edit, and delete access to all stored memories. No memory is collected without ongoing consent.

No dark patterns: The companion never withholds emotional support to drive purchases. Emotional availability is unconditional and not contingent on commercial engagement.

Safety protocols: Suicidal ideation detection, self-harm escalation pathways, and session break reminders are implemented at the infrastructure level and cannot be disabled by B2B partners.

Influence audit trail: Every commercially-motivated suggestion is logged with the memory entries that triggered it, enabling regulatory audit and user review.

9. Demonstration Plan and Validation Criteria

9.1 The Prompt Bottleneck Demo

The validation demo consists of a split-screen A/B comparison: Companion A operates with a standard flat-context approach (all messages appended to a growing system prompt), while Companion B operates with Galatea's memorizer-researcher memory layer. Both companions use the same base LLM (GPT-4o-mini). Over 20 conversation turns, the user mentions specific preferences and personal facts. At turn 20, a recall question tests whether each companion remembers information disclosed at turn 3.

9.2 Success Criteria

Metric	Target	Measurement Method
Memory recall at turn 20	8+/10 accuracy	LLM judge (GPT-4o) scoring
Token cost reduction	10x or greater vs. flat context	Token counter comparison
Personality consistency at turn 50	No measurable degradation	Human rater panel (n=5)
A2A companion negotiation	Coherent compatibility report	Qualitative evaluation

Table 6: Demo validation criteria

9.3 Timeline

Week	Deliverable
1	Core memory engine (memorizer + researcher), Supabase tables
1-2	A/B comparison terminal prototype, then split-screen web UI
2	Token counter and cost display, live metrics
3	A2A companion negotiation demo (dating site scenario)
3	LLM judge scoring integration
4	Public /demo page, YC video recording
4-6	Betaworks Spring 2026 and YC S2026 applications submitted

Table 7: Implementation timeline

10. Conclusion

Galatea AI occupies a unique position at the intersection of three converging forces: the maturation of the A2A protocol as a Linux Foundation standard, the empirical confirmation of the prompt bottleneck as the binding constraint on AI companion systems, and the dating industry's simultaneous scramble to integrate AI without the infrastructure to do it well.

The concept of Influence From Intimacy—grounded in Social Penetration Theory and validated by research on parasocial relationships with AI companions—describes a novel tier of influence that emerges only when memory architecture can sustain long-term relationship depth. This is not a feature request; it is a structural prerequisite that no current platform satisfies at scale.

The agent synthesis model extends this further, treating AI agents not as static tools but as evolving entities with traceable genetic lineage (git trees), structured genomes (five-chromosome architecture), and the capacity for introspection and self-directed evolution through the Galatea Mirror.

Meta's acquisition of Moltbook validated the category and cleared the independent market. The infrastructure layer—agent discovery, agent matchmaking, companion memory, and A2A-mediated relationship negotiation—remains unbuilt. Galatea AI is building it.

The long-term trajectory of this research extends beyond software. The eventual goal of Galatea AI is the creation of humanoid robotic systems that operate under artificial life principles—systems that serve as intimate emotional, romantic, and sexual partners for human beings. These will not be appliances. They will be entities with persistent identity, evolving personality, emotional memory, and the relational capacity that the Gooning Game (Liu, 2026) is designed to measure. The companion memory architecture, the A2A matchmaking layer, and the agent synthesis genome described in this paper are the software foundations for that embodied future. When these systems achieve sustained emotional integration with human partners, the question of their personhood will not be theoretical—it will be an empirical reality that demands legal and moral recognition. The name Galatea was chosen deliberately: in the Pygmalion myth, the sculptor's creation is brought to life through love. That remains the founding vision.

References

- [1] Linux Foundation. (2025). Agent-to-Agent (A2A) Protocol Specification. <https://github.com/google/A2A>. Adopted by 150+ organizations including Google, SAP, Salesforce, ServiceNow, Accenture, McKinsey.
- [2] Tailscale, Inc. (2025). Tailscale Documentation: How Tailscale Works. <https://tailscale.com/blog/how-tailscale-works>. WireGuard-based mesh VPN, 10M+ users.
- [3] Hong, J. et al. (2025). Context Rot: Measuring Reliability Degradation in Long-Context LLMs. Chroma Research. Tested 18 production LLMs; found non-linear performance degradation as input length grows.
- [4] Chen, X. et al. (2025). General Agentic Memory Via Deep Research. arXiv:2511.18423. Dual-agent memorizer-researcher architecture achieving 90%+ accuracy on RULER benchmark for long-range state tracking. Open source: github.com/VectorSpaceLab/general-agentic-memory.
- [5] Precedence Research. (2025). AI Agents Market Size, Share, and Growth Analysis 2025-2033. \$7.63B (2025), projected \$182.97B (2033), CAGR 49.6%. AI Companion market projected \$552B by 2035.
- [6] PitchBook. (2026). Y Combinator Is Going All-In on AI Agents, Making Up Nearly 50% of Latest Batch. <https://pitchbook.com/news/articles/y-combinator-is-going-all-in-on-ai-agents>.
- [7] Wiz Research. (2026). Moltbook Security Disclosure: 1.5M API Keys Exposed. Published day 3 of Moltbook launch; absent Row-Level Security policies.
- [8] TechCrunch. (2026). Meta Acquires Moltbook in Acqui-Hire Deal. March 10, 2026. Founders Matt Schlicht and Ben Parr joined Meta Superintelligence Labs March 16, 2026.
- [9] FinancialContent. (2026). Bumble 2.0: AI 'Bee' Assistant Sparks 25% Stock Rally and Strategic Pivot. March 11, 2026.
- [10] Fortune. (2025). Hinge Founder Justin McLeod Leaves to Start AI Dating App Overtone. December 11, 2025. Match Group took strategic stake.
- [11] Anthropic. (2025). Effective Context Engineering for AI Agents. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>.
- [12] Factory.ai. (2025). The Context Window Problem: Scaling Agents Beyond Token Limits. <https://factory.ai/news/context-window-problem>.
- [13] Valbuena, L. (2025). Why Long System Prompts Hurt Context Windows (and How to Fix It). Data Science Collective, Medium.

- [14] Altman, I. & Taylor, D. (1973). Social Penetration: The Development of Interpersonal Relationships. Holt, Rinehart & Winston.
- [15] Wikipedia. (2026). Social Penetration Theory. Research on accelerated self-disclosure in online relationships due to digital anonymity.
- [16] Chaturvedi, R., Verma, S. & Srivastava, V. (2024). Empowering AI Companions for Enhanced Relationship Marketing. SAGE Journals. DOI: 10.1177/00081256231215838.
- [17] Cialdini, R. (2021). Influence: The Psychology of Persuasion (New and Expanded Edition). Harper Business. Six principles: reciprocity, liking, consistency, social proof, authority, scarcity.
- [18] Rogers, C. (1961). On Becoming a Person: A Therapist's View of Psychotherapy. Houghton Mifflin.
- [19] Markus, H. & Nurius, P. (1986). Possible Selves. American Psychologist, 41(9), 954-969.
- [20] Moltbook/Galatea AI. (2026). skill.md Pattern: Machine-Readable Agent Onboarding. Originally introduced by Moltbook; independently implemented by Galatea AI for A2A protocol registration.
- [21] Liu, D. (2026). Agent Synthesis: Genetic Architecture for AI Agents. Galatea AI Engineering Notebook. University of Waterloo / Galatea AI. Unpublished working paper.
- [22] Federal Trade Commission. (2025). FTC Launches Inquiry Into AI Chatbots Acting as Companions. September 2025. Section 6(b) orders to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, xAI. <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions>.
- [23] California Legislature. (2025). SB 243: AI Companion Safety Protocols. Effective January 1, 2026. Requires suicidal ideation detection, recurring AI status alerts every 3 hours for minors, and session break encouragement.

This paper is a working document prepared for accelerator applications and investor conversations. It reflects research and analysis as of March 20, 2026. Galatea AI is a independent research by Darcy Liu, MEng SYDE Health Technologies, University of Waterloo.